


# Cloud Native Observability

---

With Prometheus, ElasticStack and Jaeger

# whoami

- Alexander Knipping, Systems Engineer @ noris network AG
  - Kubernetes, Prometheus
  - Cloud Native Adoption
-  Open Source
  - <https://github.com/obitech>
- Co-Host of CNCF Nürnberg Meetup
  - <https://www.meetup.com/Kubernetes-Nurnberg/>
  - <https://github.com/k8s-nue-meetup/talks>

# Agenda

## 1. Problem Space

- a. Why are we talking about this?
- b. What is Observability?

## 2. Tools & Demo

- a. Demo Microservice Application
- b. Prometheus, ElasticStack, Jaeger


## 3. Q&A, Feedback



# Why are we talking about this?


---

# The Traditional Architecture




# The Traditional Architecture

- Easy to develop, deploy, test, monitor
- Scaling becomes inefficient
  - Up vs. out
  - Only scales in regards to transactions
  - Only parts need to scale
- Code base grows
  - Introduces tight coupling, side effects
  - Onboarding new developers becomes harder
  - Deployments are tricky
- Big machines are expensive




# Microservices & Distributed Systems



# Microservices & Distributed Systems

- Small services
  - Easy to understand & reason about
  - Faster testing, building, deploying
  - Loosely coupled
- Higher isolation
  - Single degraded service is unlikely to bring down whole system
  - Independent deployments
- DevOps required
- Increased infrastructure complexity
  - Distributed System
  - Deployment complexity
  - Losing visibility/**observability**





# What is Observability?

- What is going on in my system?
- Monitoring
  - Resource metrics (CPU, RAM, Disk, Net, ...)
  - Latency distributions
  - Error rates
  - Application & business metrics
- Logging
  - Syslog streams/files
  - Application Logs
- Tracing
  - Which service is talking with whom and for how long?
  - Communication graphing


# Tools & Demo

---

# Demo Microservice Application

<https://github.com/obitech/micro-obs>

- Two services, REST APIs
- User → Item
- User → Order → Item → Order
- Instrumented directly (no sidecar)
- Written in Go



# Cloud Native Observability

## ... using Open Source tools:

### Monitoring with Prometheus

Resource utilization, error rate,  
latency, etc.



### Centralised Logging with ElasticStack

Gather & query service logs

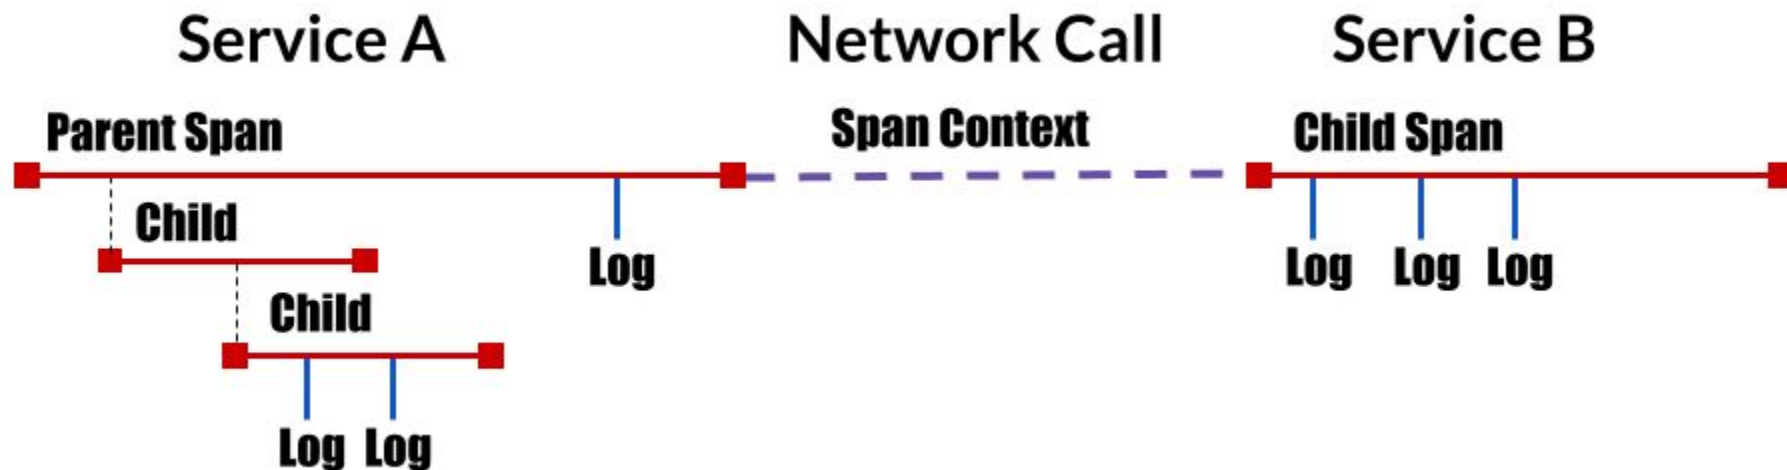


### Distributed Tracing with OpenTracing & Jaeger

Visualize service  
communication

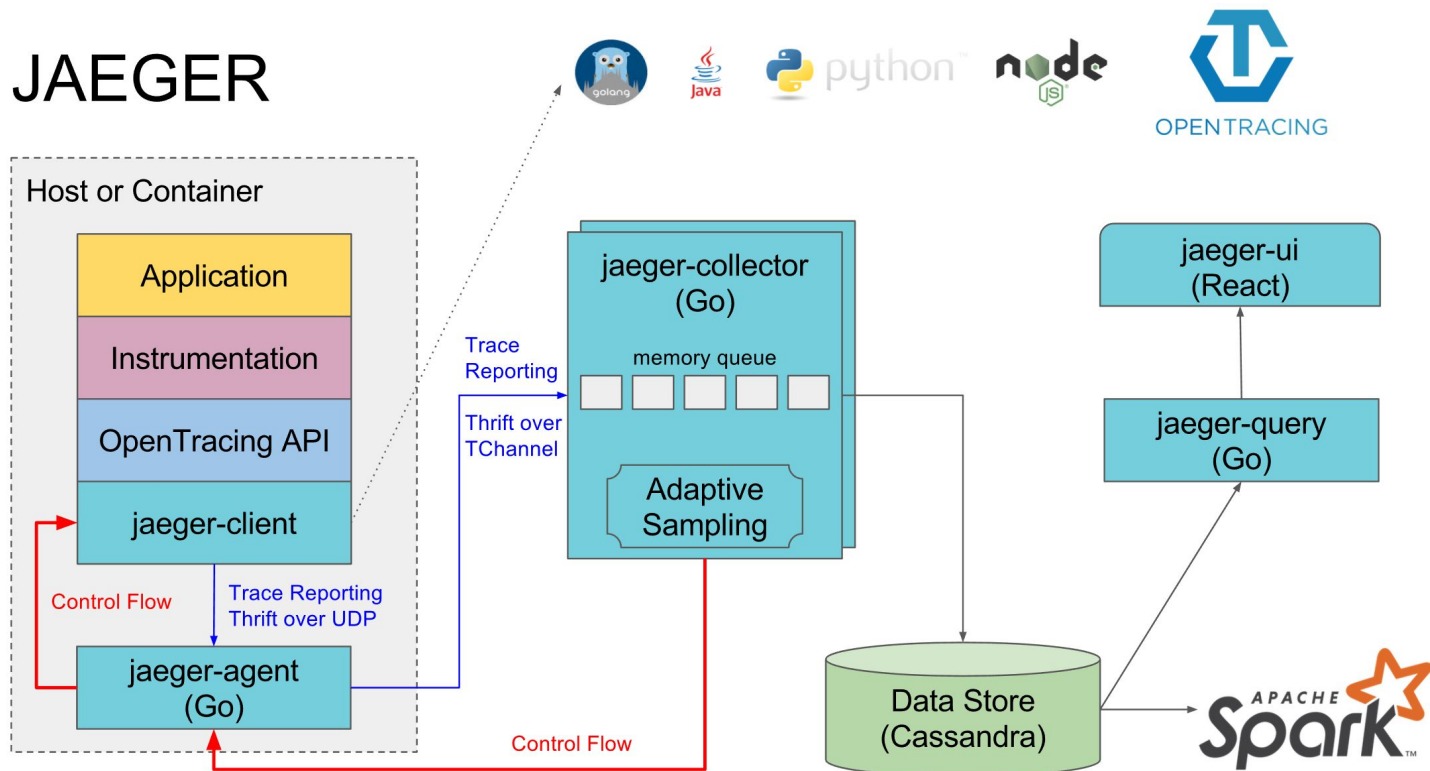


# OpenTracing / OpenTelemetry




<https://opentracing.io/docs/overview/>

# JAEGER



# Demo: Jaeger

# ElasticStack

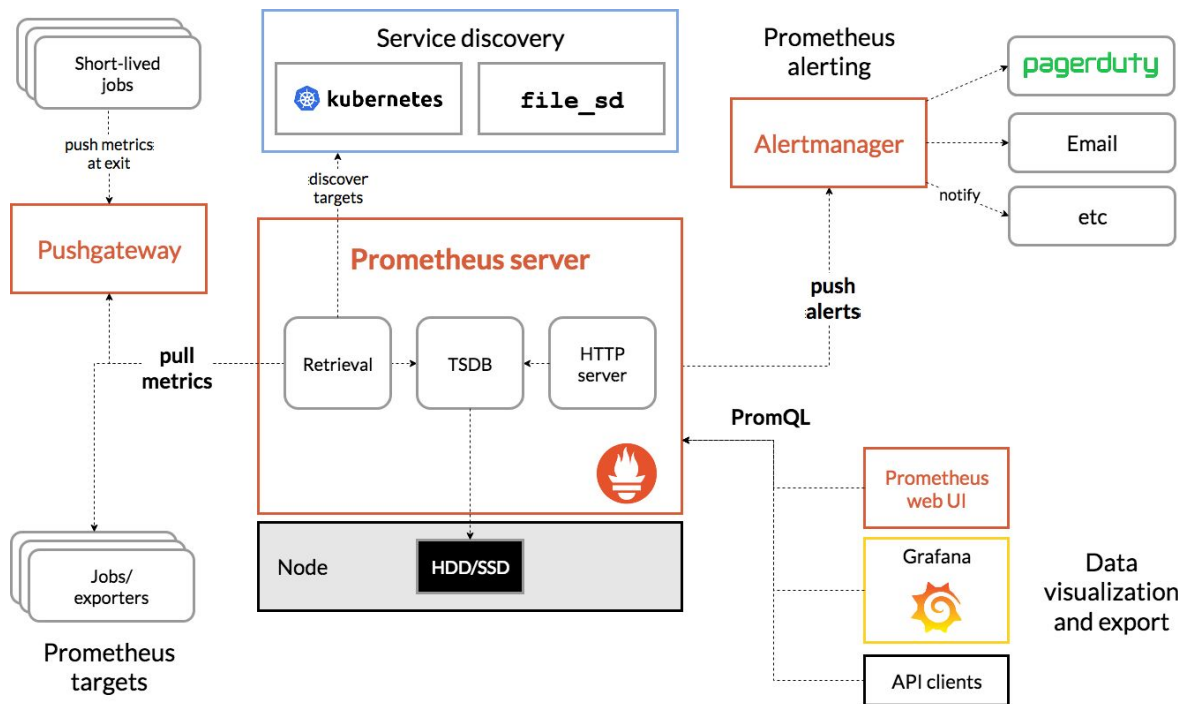


<https://www.elastic.co/guide/en/beats/libbeat/current/beats-reference.html>



# Demo: ElasticStack

# Prometheus



<https://prometheus.io/docs/introduction/overview/>

# Demo: Prometheus

Demo: Show some code?

# Final thoughts

- Follow best practices
  - Structured logs
  - Log request IDs, not trace IDs
  - Monitoring...
    - Resources: USE method
    - Endpoints: RED method
- Simple instrumentation, no overhead for application
- But: complex observability stack
- Logstash & Elasticsearch = memory hungry (~50%, 1GB)

Thank you!

