


# Grafana Loki

---

How to make sense of your logs

# whoami

- Alexander Knipping, IT Systems Engineer @ noris network
  - Observability
  - Cloud Native Adoption
-  Open Source
  - Contributor to Kubernetes ([wg-component-standard](#)) and [M3DB](#)
  - <https://github.com/obitech>
- Co-Host of this Meetup

# KCD 2020

- June/July 2020 in Munich
- Would you like to be part...
  - ... as a sponsor?
  - ... as a speaker?



**Kubernetes Community Days**

— **MUNICH 2020** —

# Agenda

## Grafana Loki: “Prometheus but for Logs”

1. Overview ElasticSearch
2. Overview Prometheus
3. Overview Loki
4. Demo



# Elasticsearch

---

# Elasticsearch: Overview

- OpenCore, distributed search- and analytics engine based on Apache Lucene
- Released in 02/2009 by Elastic
- Used to save & analyse text documents
- HTTP REST API
- Often deployed as part of the Elastic- or ELK-Stack

# Elasticsearch: Inverted Index

Three documents:

1 → {"hello", "alex"}

2 → {"my", "name"}

3 → {"is", "alex"}

Inverted index:

"hello" → [1:0]

"alex" → [1:0, 3:1]


"my" → [2:0]

...

- Stemming: rainfall → rain
- Synonyms: Pen, Pens → Pen

# Elasticsearch: Sharding

- Split up index over several machines (shards)
- Logs: usually one index per day: **logstash-2019-12-11**
- Given **logstash-2019-12-11** has 150 log entries, 3 shards and replication factor 2:





# Elasticsearch: Summary

## What we get:

- Ability to index...
  - ... big amount of documents (logs)
  - ... in a fine grained way
- Ability to ask elaborate questions
  - Show me the Top 10 Users by amount paid aggregated across countries during the past two months

## We pay for it with:

- Resource intensity
- Operational complexity



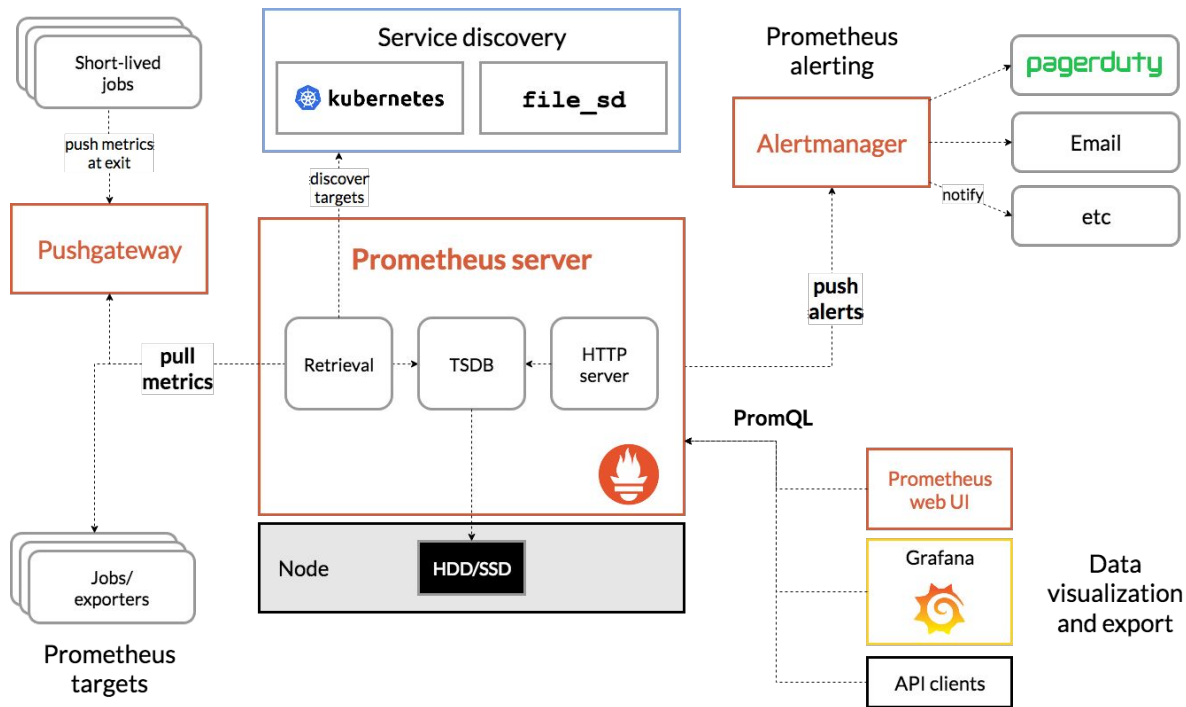
# Prometheus

---

# Prometheus: Overview

- Metric-base white-box monitoring & alerting system
- CNCF [Graduated project](#), started at SoundCloud in 2012
- Multidimensional data model
- Simple to operate
- Optimized for dynamic (cloud) environments

# Prometheus: Architecture



# Prometheus: TSDB

```
prometheus_sd_discovered_targets{config="grafana",name="scrape"} 1
prometheus_sd_discovered_targets{config="node-exporter",name="scrape"} 1
prometheus_sd_discovered_targets{config="prometheus",name="scrape"} 1
prometheus_sd_discovered_targets{config="telegraf",name="scrape"} 1
```

- config="grafana"
  - → [000001]
- config="node-exporter"
  - → [000002]
- name="scrape"
  - → [000001, 000002]

```
./data
├── b-000001
│   ├── chunks
│   │   ├── 000001
│   │   └── 000002
│   ├── index
│   └── meta.json
├── b-000002
│   ├── meta.json
│   └── wal
│       ├── 000001
│       ├── 000002
│       └── 000003
```

☐ Enable query history

```
100 - avg by (instance) (  
  (node_memory_MemAvailable_bytes * 100) /  
  node_memory_MemTotal_bytes  
)
```

Load time: 181ms  
Resolution: 14s  
Total time series: 1

Execute

- insert metric at cursor - ▾

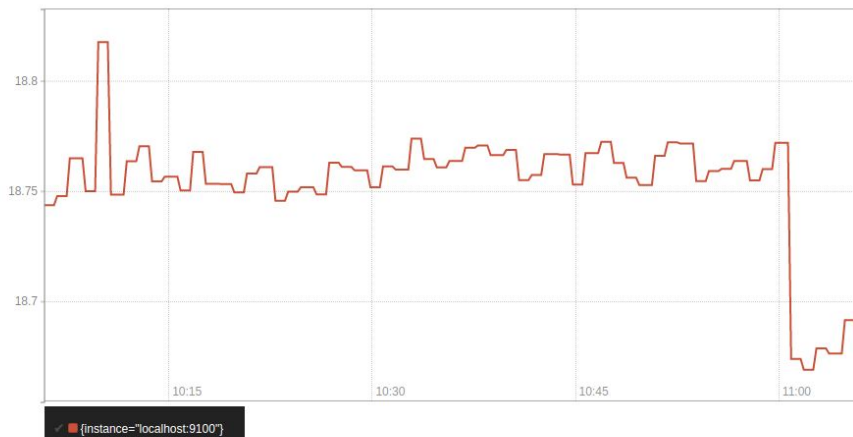
Graph

Console

- 1h +

◀ Until ▶

Res. (s)

☐ stacked[Remove Graph](#)[Add Graph](#)

# Prometheus: Cardinality Explosion

- Cardinality is multiplicative
- Total number of time series = metric names \* cardinality of labels \* targets
- Prometheus can handle a couple 100k time series

# Multiplicative Cardinality

You have a histogram, *http\_request\_duration\_seconds\_bucket*

100 instances \* 10 buckets -> 1000 series

\* 10 endpoints -> 10,000 series

\* 10 response codes -> 100,000 series (maximum recommended cardinality)

\* 4 http methods -> 400,000 series

\* 100 tenants -> 40,000,000 series ☠☠☠☠

splunk > turn data into doing

<https://promcon.io/2019-munich/slides/containing-your-cardinality.pdf>



# Prometheus: Summary

## What we get:

- Low operational cost
  - Simplicity
  - Robustness
- High throughput
  - ~1,5m samples/s
  - ~500k time series
- Multidimensional data model
- Rich query language

## We pay for it with:

- Single Server only
  - Label cardinality is limiting factor
- No built-in long term storage



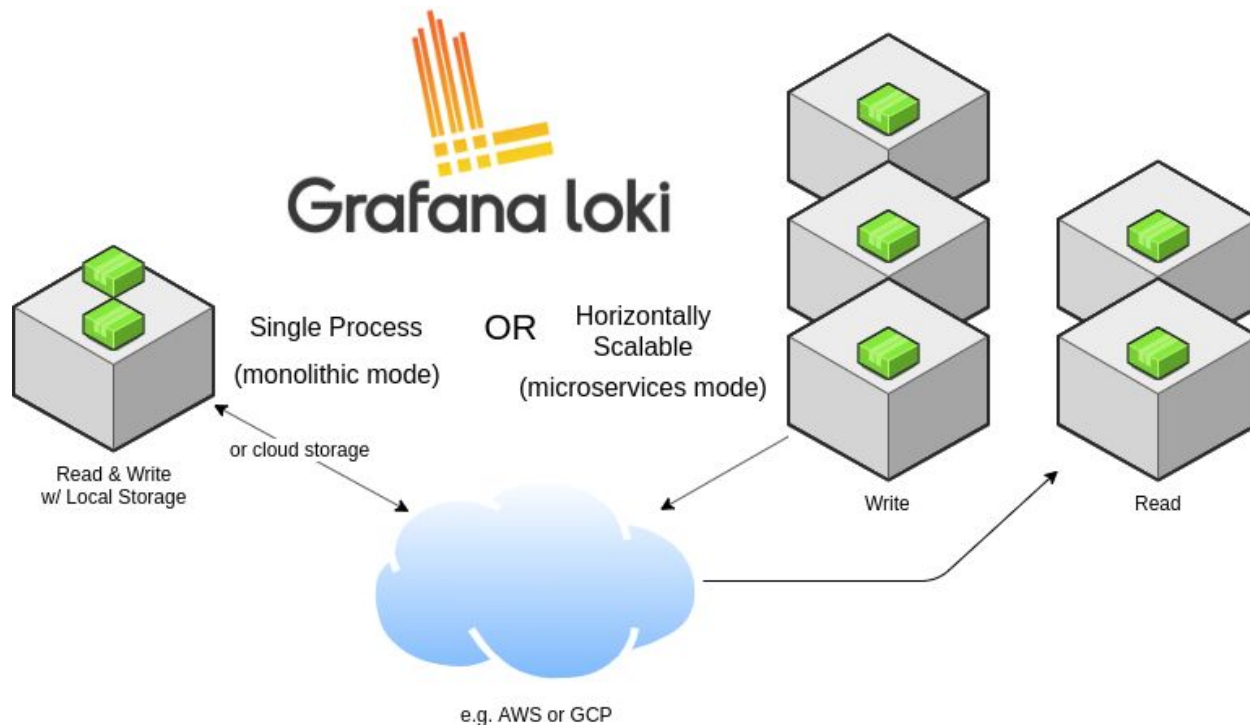
# Grafana Loki

---

# Grafana Loki: Overview

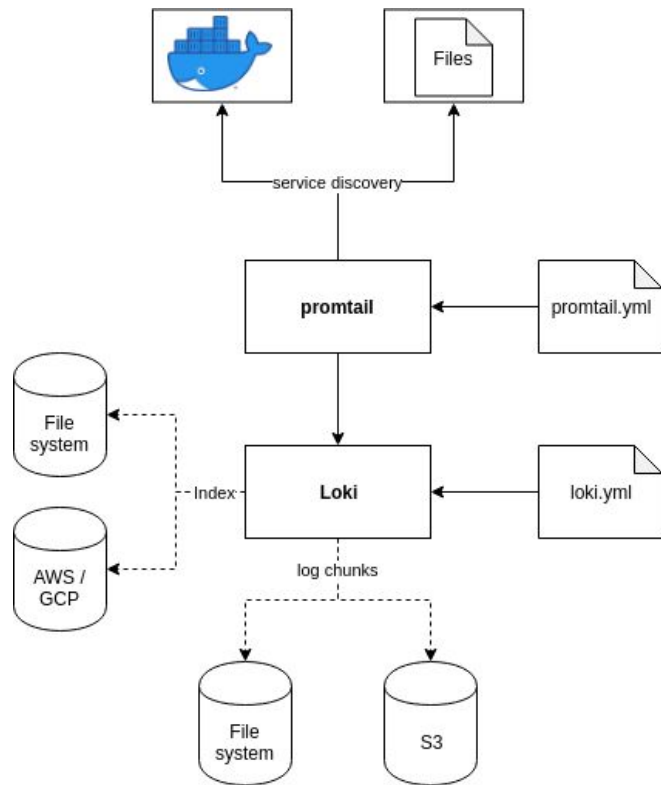
- Started 2017, v1.0 since 12/2019
- “Prometheus but for Logs”
  - Complexity of Elasticsearch often not needed for operational log aggregation
  - Index only metadata (labels) of logs
- Column store for indices, S3 for chunks
- Operationally simple
  - Single process & microservices mode
- First-class Grafana integration
  - Less context switches

# Grafana Loki: Architecture



# Grafana Loki

- Promtail discovers logs, adds labels and sends them to Loki
- Usage of Prometheus libraries for service discovery and label relabeling
- In Loki, log streams are batched up into chunks, compressed, then sent to chunk store
- Indices (labels) are sent to column store
- On read, chunks are streamed out via gRPC, deduplicated and sent out via HTTP/1



# Demo

---

<https://github.com/noris-network/loki-demo>

# Grafana Loki: Summary

## What we get:

- Simplicity
- Scalability
- Cost effectiveness
- Less context switches between metrics and logs

## We pay for it with:

- Restricted by label cardinality
- At some point we need to use cloud services...
  - ... or run our own Cassandra & S3

# Thank You!

Demo: <https://github.com/noris-network/loki-demo>

Loki: <https://github.com/grafana/loki>

Loki Design Doc: [on Google docs](#)